

AML Report part-2

by Sadikshya Duwadi

Submission date: 10-May-2024 11:14PM (UTC+0545)

Submission ID: 2376197864

File name: R_report_part-2.docx (3.54M)

Word count: 1438

Character count: 7974

“Data Analysis on Concrete Strength Part -2”

PART II

Sadikshya Duwadi
BSc (hons) Computing, 2024

Summary

The first part of the report was started by combining the two concrete strength datasets ‘test’ and ‘train’ followed by executing preliminary tasks. These tasks encompassed the following: Exploratory Data Analysis (EDA) to uncover hidden patterns and trends in the data, Data Visualization and Principal Component Analysis (PCA) to reduce dimensionality. Data pre-processing was done to find missing values and overall missingness in the data and the missing values were handled using mean imputation. Finally, the outliers were detected and removed. The respective means of the variables were used to replace any missing values and the variables were scaled. The final dataset contained 1030 rows and 10 columns.

Modelling

Equipped with the pre-processed data, five different machine learning algorithms were developed to predict the concrete strength.

Data frame of combined scaled numeric data with non-numeric data

```
> # Extracting non-numeric columns
> dataset <- final_imputed_data_without_outliers[, !sapply(final_imputed_data_without_outliers, is.numeric)]
> str(final_imputed_data_without_outliers)
'data.frame': 1030 obs. of 10 variables:
 $ Cement      : num  540 540 332 199 266 ...
 $ Blast.Furnace.Slag: num  0 0 142 132 114 ...
 $ Fly.Ash      : num  0 0 0 0 0 0 0 0 0 ...
 $ Water        : num  162 162 228 192 228 228 192 192 228 ...
 $ Superplasticizer : num  2.5 2.5 0 0 0 0 0 0 0 ...
 $ Coarse.Aggregate : num  1040 1055 932 978 932 ...
 $ Fine.Aggregate  : num  676 676 NA 826 670 ...
 $ Age          : int   28 28 NA NA 90 28 28 90 28 NA ...
 $ Strength      : num  NA 61.9 40.3 44.3 47 ...
 $ isTrain       : chr   "train" "train" "train" "train" ...
>
> # Combining scaled numeric data with non-numeric data
> final_data <- cbind(num_data, dataset)
> str(final_data)
'data.frame': 1030 obs. of 10 variables:
 $ Cement      : num  540 540 332 199 266 ...
 $ Blast.Furnace.Slag: num  0 0 142 132 114 ...
 $ Fly.Ash      : num  0 0 0 0 0 0 0 0 0 ...
 $ Water        : num  162 162 228 192 228 228 192 192 228 ...
 $ Superplasticizer : num  2.5 2.5 0 0 0 0 0 0 0 ...
 $ Coarse.Aggregate : num  1040 1055 932 978 932 ...
 $ Fine.Aggregate  : num  676 676 NA 826 670 ...
 $ Age          : int   28 28 NA NA 90 28 28 90 28 NA ...
 $ Strength      : num  NA 61.9 40.3 44.3 47 ...
 $ dataset      : chr   "train" "train" "train" "train" ...
>
> #Splitting the dataset
> strength_train<-final_data [final_data $isTrain=="train",]
> strength_test<-final_data [final_data $isTrain=="test", ]
> |
```

Figure 1.1 Structure of Combined numeric data and dataset

Data frame of test and train data after removing the column is_Train.

```
> #removing variable isTrain from both train and test
> strength_train$isTrain<-NULL
> strength_test$isTrain<-NULL
> View(strength_train)
> View(strength_test)
> str(strength_train)
'data.frame': 642 obs. of 9 variables:
 $ Cement      : num 540 266 266 199 199 ...
 $ Blast.Furnace.Slag: num 0 114 114 132 132 ...
 $ Fly.Ash      : num 0 0 0 0 0 94 0 0 0 ...
 $ Water        : num 162 228 228 192 192 228 228 192 192 ...
 $ Superplasticizer : num 2.5 0 0 0 0 0 0 0 0 ...
 $ Coarse.Aggregate : num 1055 932 932 978 978 ...
 $ Fine.Aggregate  : num 676 670 670 826 826 ...
 $ Age           : int 28 90 28 90 28 90 28 90 100 28 ...
 $ Strength       : num 61.9 47 45.9 38.1 28 ...
> str(strength_test)
'data.frame': 281 obs. of 9 variables:
 $ Cement      : num 349 140 313 425 475 ...
 $ Blast.Furnace.Slag: num 0 209 262 114 119 ...
 $ Fly.Ash      : num 0 0 0 0 0 0 0 0 0 ...
 $ Water        : num 192 192 176 151 181 ...
 $ Superplasticizer : num 0 0 8.6 18.6 8.9 12.1 16.5 11.6 10.3 15.9 ...
 $ Coarse.Aggregate : num 1047 1047 1047 936 852 ...
 $ Fine.Aggregate  : num 807 807 612 804 782 ...
 $ Age           : int 3 7 3 3 3 3 3 3 1 3 ...
 $ Strength       : num 15.1 14.6 28.8 36.3 37.8 ...
```

Figure 1.2 Structure of Train and Test Datasets after removing is_Train

1. Linear Regression

13

6

The value of an independent variable is used to predict the value of a dependent variable using linear regression analysis. A straight line or surface that minimises the differences between the expected and actual output values is fitted using linear regression.

```
> # Training the model
> model <- train(formula, data = strength_train, method = "lm")
> summary(model)

Call:
lm(formula = .outcome ~ ., data = dat)

Residuals:
    Min       1Q   Median       3Q      Max
-25.6247  -4.8739  -0.2046   5.2629  30.3648

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  33.879180   26.309307   1.288  0.19831
Cement        0.102604    0.008149  12.591 < 2e-16 ***
Blast.Furnace.Slag  0.079614    0.009785   8.136 2.16e-15 ***
Fly.Ash        0.047584    0.012100   3.933 9.33e-05 ***
Water         -0.228637    0.040834  -5.599 3.21e-08 ***
Superplasticizer  0.343216    0.104494   3.285  0.00108 **
Coarse.Aggregate -0.001508    0.009262  -0.163  0.87070
Fine.Aggregate  -0.007078    0.010872  -0.651  0.51524
Age            0.319834    0.011453  27.925 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.145 on 633 degrees of freedom
Multiple R-squared:  0.766,    Adjusted R-squared:  0.763
F-statistic: 259 on 8 and 633 DF, p-value: < 2.2e-16
```

Figure 1.1.1 Summary

```
> # Predicting on the test set
> predictions <- predict(model, newdata = strength_test)
> predictions
```

731	733	737	738	739	740	742	743	744	745	746	747
19.458509	15.923932	44.775750	52.188975	47.865245	42.654671	49.178813	46.594405	45.345685	50.091262	41.052972	47.377286
748	749	750	752	753	754	755	756	757	758	760	761
46.192371	51.190600	52.855279	43.934006	51.190600	47.873740	50.190654	46.981377	33.723872	57.907107	55.861086	57.907107
762	763	764	765	767	768	769	770	771	772	773	775
54.255013	58.087103	63.049134	66.862449	63.545588	74.260222	74.826780	77.256396	78.213759	75.425678	74.739766	69.198332
776	777	778	779	780	781	782	783	784	785	786	787
12.903864	43.927727	15.718257	20.195928	29.151270	43.223950	33.330849	48.249487	18.652486	17.543091	22.899354	27.377025
788	789	790	791	793	794	795	796	797	798	799	800
36.332367	13.723426	30.674609	46.600657	57.976886	31.697941	45.605699	16.338742	32.985562	24.541399	20.067693	28.063534
801	802	803	804	805	806	807	809	810	811	812	813
28.894538	37.594357	21.556955	18.518760	29.552796	17.077868	27.817356	26.127535	27.688858	21.690671	29.686512	38.641854
814	815	816	817	818	819	820	821	822	823	824	825
58.593047	34.044761	45.344926	33.557744	56.769799	45.983342	37.064498	29.475013	60.778023	34.823430	44.293596	58.366277
826	827	828	829	830	831	832	833	834	835	836	837
37.965452	40.034630	41.241348	64.152151	32.123606	34.750091	49.074789	28.915265	31.154101	47.628385	23.589850	40.541033
838	839	840	841	842	843	844	845	846	847	848	849
56.005340	54.668320	51.617591	42.123532	25.011826	47.284179	50.954843	22.378146	24.945546	19.791282	24.167812	24.845899
850	851	852	853	854	855	856	857	858	859	860	861
17.235973	22.550186	24.959833	29.512708	27.787123	32.843155	23.507130	33.608590	23.365465	34.187156	36.742465	41.798497
862	863	864	865	866	867	868	869	870	871	872	873
42.563932	32.320807	54.739648	50.815145	56.935148	52.663938	48.047827	55.532706	47.993106	63.218773	42.890626	58.252181
874	875	876	877	878	879	880	881	882	883	884	885

Figure 1.1.2 Predictions of Linear Regression

Above Figure 1.1.1 shows the summary of the linear regression model and Figure 1.1.2 shows the prediction made by linear regression model in test dataset.

```
> # Printing the metrics
> print(paste("R2: ", r2))
[1] "R2: 0.747564646718755"
> print(paste("Adjusted R2: ", adj_r2))
[1] "Adjusted R2: 0.740140077504601"
> print(paste("MSE: ", mse))
[1] "MSE: 69.8687558871473"
> print(paste("RMSE: ", rmse))
[1] "RMSE: 8.35875324956702"
> print(paste("MAE: ", mae))
[1] "MAE: 6.48433978334516"
> # Creating a data frame with the actual and predicted values
> comparison <- data.frame(Actual = strength_test$Strength, Predicted = predictions)
> # Creating a scatter plot
> ggplot(comparison, aes(x = Actual, y = Predicted)) +
+   geom_point() +
+   geom_smooth(method = lm, se = FALSE, color = "blue") +
+   labs(title = "Actual vs Predicted", x = "Actual", y = "Predicted") +
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'
```

Figure 1.1.3 Evaluation Metrics

The value for evaluating the metrics is printed in the above fig.1.1.3. Different performance metrics: R-squared (R2), Adjusted R-squared (Adj. R2), Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) values are obtained for linear regression model.




Figure 1.1.4 Scatter Plot for Linear Regression with actual vs predicted values

The predicted values of a linear regression model is shown against the actual values in the above scatter plot.

2. Random Forest

A popular machine learning approach called Random Forest mixes the outputs of several decision trees to get a single outcome. Its versatility and ease of use, combined with its ability to handle both regression and classification issues, makes it more flexible to use.

```

> # Training the model
> model <- randomForest(formula, data = strength_train)
> plot(model)
> # Printing the model summary
> print(summary(model))

```

	Length	Class	Mode
call	3	-none-	call
type	1	-none-	character
predicted	642	-none-	numeric
mse	500	-none-	numeric
rsq	500	-none-	numeric
oob.times	642	-none-	numeric
importance	8	-none-	numeric
importanceSD	0	-none-	NULL
localImportance	0	-none-	NULL
proximity	0	-none-	NULL
ntree	1	-none-	numeric
mtry	1	-none-	numeric
forest	11	-none-	list
coefs	0	-none-	NULL
y	642	-none-	numeric
test	0	-none-	NULL
inbag	0	-none-	NULL
terms	3	terms	call

Figure 1.2.1 Summary of Random Forest




Figure 1.2.2 Plotting the Model value

The model value after using randomForest() on Strength_Train data is plotted and is shown in fig.1.2.2, where the plot of the errors in y-axis is from 0-110 as the number of trees rises from 0 to 500 in x-axis using a random forest model trained on train data.

```

> # Printing the metrics
> print(paste("R2: ", r2))
[1] "R2: 0.875868771262608"
> print(paste("Adjusted R2: ", adj_r2))
[1] "Adjusted R2: 0.876310519407581"
> print(paste("MSE: ", mse))
[1] "MSE: 37.3859024943196"
> print(paste("RMSE: ", rmse))
[1] "RMSE: 6.11440123759634"
> print(paste("MAE: ", mae))
[1] "MAE: 4.50949076974893"

```

Figure 1.2.3 Evaluation Metrics

The value for evaluating the metrics is printed in the above fig.1.2.3. Different performance metrics: R-squared (R2), Adjusted R-squared (Adj. R2), Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) values are obtained for Random Forest model.

```

> # Creating a data frame with the actual and predicted values
> comparison <- data.frame(Actual = strength_test$Strength, Predicted = predictions)
> # Creating a scatter plot
> ggplot(comparison, aes(x = Actual, y = Predicted)) +
+   geom_point() +
+   geom_smooth(method = lm, se = FALSE, color = "blue") +
+   labs(title = "Actual vs Predicted", x = "Actual", y = "Predicted") +
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'

```

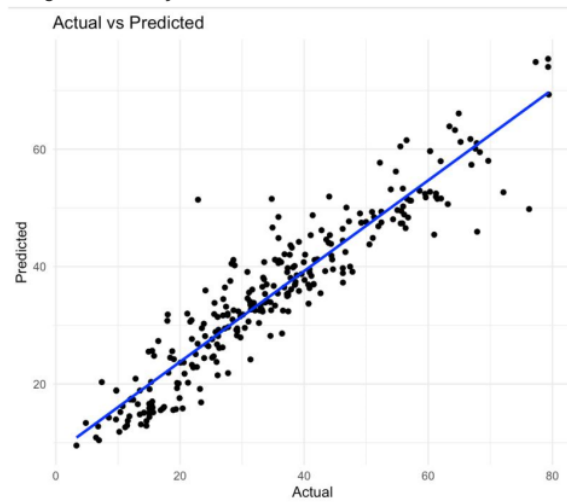
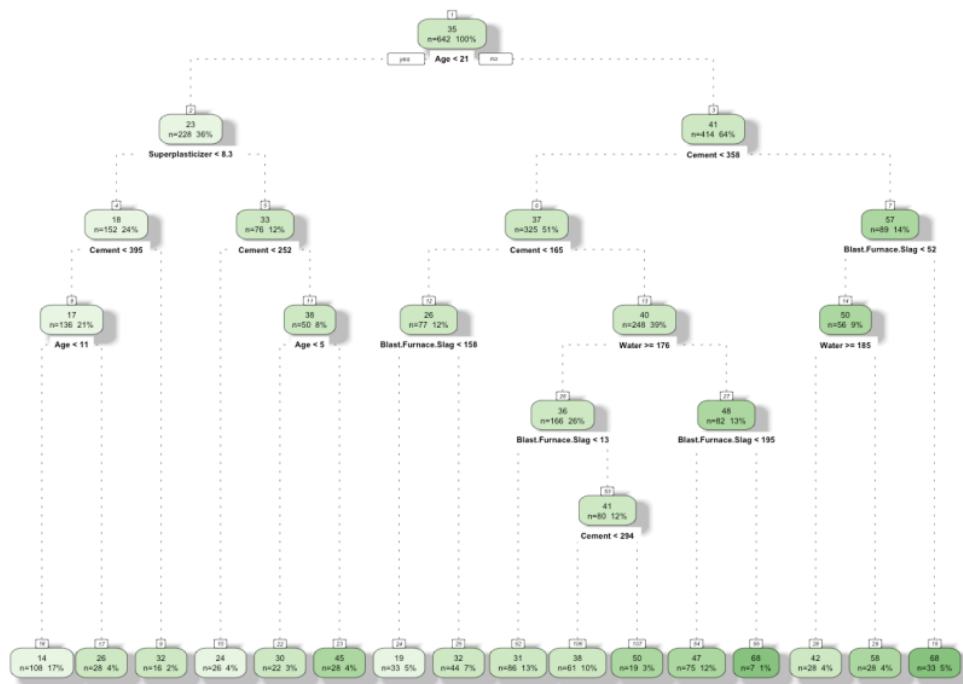


Figure 1.2.4 Scatter Plot for Random Forest with actual vs predicted values

The predicted of a *Random Forest* model is shown against the actual values in the above scatter plot. Here we can see the data seem to be more closely packed around the diagonal line than in Linear Regression Model suggesting that the *Random Forest* model would offer a much better fit to the data.

```
> # Plot the decision tree
> fancyRpartPlot(model, sub = '')
> # Plot the decision tree
> fancyRpartPlot(model, sub = '')
>
```



In fig. 1.3.1, we can see a decision tree visualisation produced by using `fancyRpartPlot()` function to forecast the target variable. The tree divides the data into multiple categories trying to predict the compressive strength of concrete. Initially, the decision tree divides the data into two categories according to whether the concrete is older than 11 years. Next, it divides the data once again according to whether or not there is less than 158 blast furnace slag. The procedure carries on until the decision tree reaches a leaf node, which has a forecast regarding the concrete's compressive strength.


```

> # Predicting on the test set
> predictions <- predict(model, newdata = strength_test)
> predictions
  731    733    737    738    739    740    742    743    744    745
14.18509 14.18509 29.66955 29.66955 29.66955 29.66955 29.66955 29.66955 29.66955 29.66955
  746    747    748    749    750    752    753    754    755    756
29.66955 44.58107 44.58107 44.58107 44.58107 44.58107 44.58107 44.58107 44.58107 44.58107
  757    758    760    761    762    763    764    765    767    768
44.58107 68.10303 68.10303 68.10303 67.84143 68.10303 68.10303 68.10303 68.10303 68.10303
  769    770    771    772    773    775    776    777    778    779
68.10303 68.10303 68.10303 68.10303 68.10303 46.66040 14.18509 30.57547 25.88750 30.57547
  780    781    782    783    784    785    786    787    788    789
30.57547 30.57547 46.66040 46.66040 14.18509 25.88750 24.15615 46.66040 46.66040 14.18509
  790    791    793    794    795    796    797    798    799    800
30.57547 30.57547 38.21016 30.57547 30.57547 14.18509 30.57547 30.57547 24.15615 46.66040
  801    802    803    804    805    806    807    809    810    811
46.66040 46.66040 14.18509 25.88750 46.66040 14.18509 46.66040 46.66040 46.66040 24.15615
  812    813    814    815    816    817    818    819    820    821
46.66040 46.66040 46.66040 46.66040 46.66040 46.66040 46.66040 46.66040 46.66040 24.15615

```

Figure 1.3.2 Predictions of test dataset.

```

> # Printing the metrics
> print(paste("R2: ", r2))
[1] "R2: 0.694009737825136"
> print(paste("Adjusted R2: ", adj_r2))
[1] "Adjusted R2: 0.695098671142484"
> print(paste("MSE: ", mse))
[1] "MSE: 84.6328556995497"
> print(paste("RMSE: ", rmse))
[1] "RMSE: 9.19961171460784"
> print(paste("MAE: ", mae))
[1] "MAE: 7.29603865850811"
>

```

Figure 1.3.3 Evaluation Metrics

In the above fig.1.3.3, the data probability is displayed for each decision: R2, Adjusted R2, MSE, RMSE, and MAE.

```

> # Creating a data frame with the actual and predicted values
> comparison <- data.frame(Actual = strength_test$Strength, Predicted = predictions)
> # Creating a scatter plot with different colors for actual and predicted values
> ggplot(comparison, aes(x = Actual, y = Predicted)) + # Initialize the plot with data
+   geom_point() + # Add a layer of points
+   geom_smooth(method = lm, se = FALSE, color = "blue") + # Add a layer of a linear regression line
+   labs(title = "Actual vs Predicted", x = "Actual", y = "Predicted") + # Add labels
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'

```

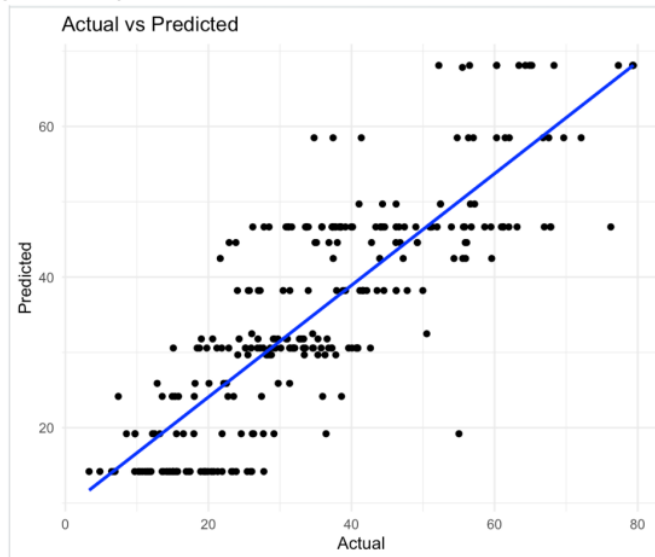


Figure 1.3.4 Scatter Plot for Decision Tree with actual vs predicted values

In the fig.1.3.4, scatter plot for Decision Tree with actual vs predicted values is shown to evaluate the model's accuracy and decision-making abilities.

4. k-Nearest Neighbours (KNN)

KNN is a member of the supervised learning domain. It is simple to grasp, nonparametric as it does not make any underlying assumptions about the distribution of the data and lazy. However, for huge datasets, it may be computationally costly.

```

> # Defining training control
> train_control <- trainControl(method = "cv", number = 10)
> model <- train(formula, data = strength_train, method = "knn", trControl = train_control)
> # Printing the model summary
> print(summary(model))
      Length Class      Mode
learn      2    -none-    list
k           1    -none-    numeric
theDots     0    -none-    list
xNames      8    -none-    character
problemType 1    -none-    character
tuneValue   1    data.frame list
obsLevels   1    -none-    logical
param       0    -none-    list
> # Predicting on the test set
> predictions <- predict(model, newdata = strength_test)
> predictions
[1] 18.32200 17.78600 49.76400 58.08000 61.47800 44.38000 50.81500 43.54000 50.50800 51.77200 42.34000
[12] 42.01667 49.76400 50.81500 58.08000 44.38000 50.81500 43.54000 62.27800 43.54000 25.07600 50.81500
[23] 61.47800 50.81500 47.43800 51.77200 72.85800 60.65800 72.61800 75.54000 76.96667 63.32000 54.37600
[34] 76.96667 76.96667 74.52714 19.69200 38.81400 19.89000 23.17600 32.50800 38.81400 32.36400 39.09800
[45] 21.57500 18.94600 23.38000 26.46000 35.82200 18.95800 29.91200 49.42600 47.94800 30.83400 40.51000
[56] 19.89000 32.66800 23.86400 20.43600 29.33400 29.19200 32.84600 17.70600 21.57500 29.61800 19.01800
[67] 25.55200 23.56200 30.93800 14.45000 27.69167 36.44400 47.23200 32.71600 41.23000 31.74200 45.73800
[78] 45.70200 31.74200 35.55200 49.38000 34.50800 36.76800 44.81200 31.71400 40.07400 36.42000 47.94800
[89] 28.52200 28.52200 48.97800 34.22400 34.22400 52.52000 19.53800 40.92200 55.88800 49.36600 51.39400
[100] 43.53200 26.36800 45.52200 48.08800 34.86800 34.86800 17.57600 31.99000 31.27400 21.08600 20.27200
[111] 22.84600 40.78200 20.27200 40.74200 22.79800 31.71400 28.92800 33.49200 25.85500 45.97200 39.24600
[122] 32.36400 51.53600 39.98400 53.37800 46.30200 42.71800 48.02800 35.00800 55.37400 40.73600 41.15000

```

Figure 1.4.1 Summary and Predictions

Above fig.1.4.1 prints the summary of the KNN model and gives the prediction on the test set data.

```

> # Printing the metrics
> print(paste("R2: ", r2))
[1] "R2: 0.681509866180901"
> print(paste("Adjusted R2: ", adj_r2))
[1] "Adjusted R2: 0.682643283027232"
> print(paste("MSE: ", mse))
[1] "MSE: 87.6233608098789"
> print(paste("RMSE: ", rmse))
[1] "RMSE: 9.36073505713514"
> print(paste("MAE: ", mae))
[1] "MAE: 7.05612862226741"
>

```

Figure 1.4.2 Evaluation Metrics on

The value for evaluating the metrics is printed in the above fig.1.2.3. Different performance metrics: R-squared (R2), Adjusted R-squared (Adj. R2), Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) values are obtained for KNN model.

```

> # Creating a data frame with the actual and predicted values
> comparison <- data.frame(Actual = strength_test$Strength, Predicted = predictions)
> # Creating a scatter plot
> ggplot(comparison, aes(x = Actual, y = Predicted)) +
+   geom_point() +
+   geom_smooth(method = lm, se = FALSE, color = "green") +
+   labs(title = "Actual vs Predicted", x = "Actual Strength", y = "Predicted Strength") +
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'

```




Figure 1.4.3 Scatter Plot for KNN with actual vs predicted strength

The above scatter plot showing the relationship between actual strength on the x-axis and predicted strength on the y-axis.

5. KNN Model with optimal value of k = 4

```

> # Printing the metrics
> print(paste("R2: ", r2))
[1] "R2: 0.679303876342072"
> print(paste("Adjusted R2: ", adj_r2))
[1] "Adjusted R2: 0.68044514368605"
> print(paste("MSE: ", mse))
[1] "MSE: 88.0285416568802"
> print(paste("RMSE: ", rmse))
[1] "RMSE: 9.38235267173859"
> print(paste("MAE: ", mae))
[1] "MAE: 7.0849709371293"
>

```

Figure 1.5.1 Evaluation Metrics

```

> # Creating a data frame with the actual and predicted values
> comparison <- data.frame(Actual = strength_test$Strength, Predicted = predictions)
> # Creating a scatter plot
> ggplot(comparison, aes(x = Actual, y = Predicted)) +
+   geom_point() +
+   geom_smooth(method = lm, se = FALSE, color = "green") +
+   labs(title = "Actual vs Predicted", x = "Actual Strength", y = "Predicted Strength") +
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'

```

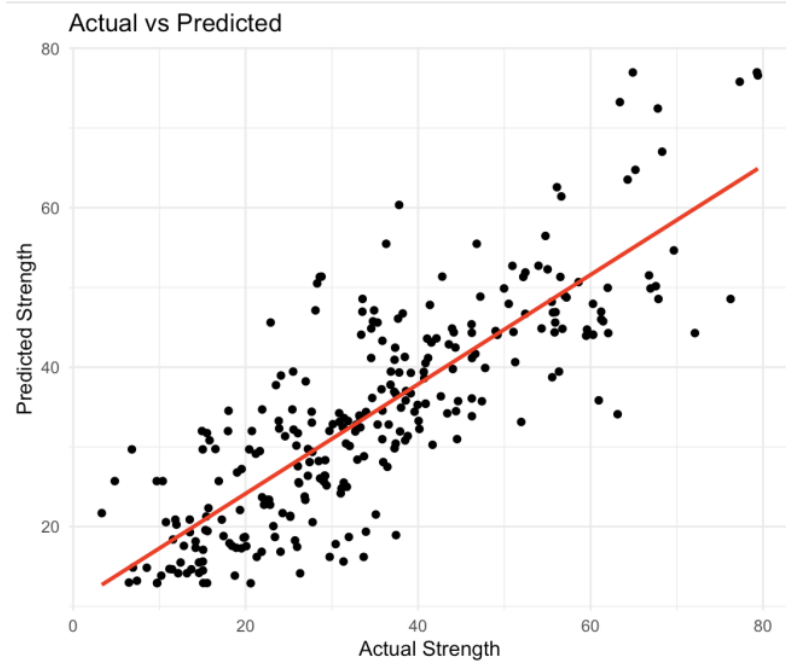


Figure 1.5.2 Scatter Plot for KNN=4 with actual vs predicted values

The above scatter plot showing the relationship between actual strength on the x-axis and predicted strength) on the y-axis with different KNN value i.e. KNN=4.

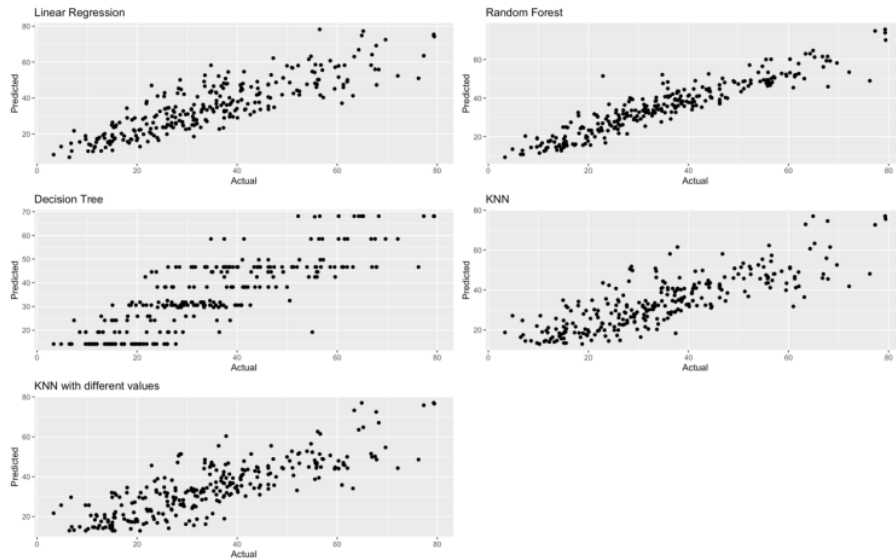


Figure 1.6 Scatter Plot for all five Models

The scatter plot for all five different models that were developed is shown to compare in the above fig.1.6. The scatter plot of the Random Forest and Linear Regression seems to be much more compacted and fit for the data than other models.

```
> # Print the metrics
> print(metrics)
      Model      R2    Adj_R2    MSE    RMSE    MAE
1 Linear Regression 0.7475646 0.7466599 69.86876 8.358753 6.484340
2 Random Forest 0.8745155 0.8740657 37.90529 6.156728 4.529004
3 Decision Tree 0.6940097 0.6929130 84.63286 9.199612 7.296039
4 KNN 0.6815099 0.6803683 87.62336 9.360735 7.056129
5 KNN (k=4) 0.6793039 0.6781544 88.02854 9.382353 7.084971
> # Select the best model based on the evaluation metrics
> best_model <- metrics[which.min(metrics$RMSE), ]
> print("Best Predicting Model based on RMSE:")
[1] "Best Predicting Model based on RMSE:"
> print(best_model)
      Model      R2    Adj_R2    MSE    RMSE    MAE
2 Random Forest 0.8745155 0.8740657 37.90529 6.156728 4.529004
>
```

Figure 1.7 Finding the Best Model

Above, the metrics for all five models are printed and the best predicting model based on evaluation metrics is shown in fig.1.7.

Model	R2	Adj. R2	MSE	RMSE	MAE
LR	0.7474646	0.7401400	69.8687558	8.3587532	6.4843397
RF	0.8758687	0.8763105	37.3859024	6.1144012	4.5094907
DECISION	0.6940097	0.6950986	84.6328556	9.1996117	7.2960386
KNN	0.6815098	0.6826432	87.6233608	9.3607350	7.0561286
KNN=4	0.6793038	0.6804451	88.0285416	9.3823526	7.0849709

An overview of the many regression models assessed using a range of performance metrics is given in this table which are: R-squared (R2), Adjusted R-squared (Adj. R2), Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE). The five models which were assessed are: Linear Regression Model, Random Forest Model, k-Nearest Neighbours Model (KNN), Decision Tree Model, and a Modified KNN model (KNN=4). We may compare and choose the best model for the task at hand by observing each model's performance across these measures based on this summary.

Model Interpretation

We can observe that, in comparison to the other models, the Random Forest (RF) model has the highest R-squared (0.8758687) and the lowest values for MSE, RMSE, and MAE indicating a superior fit to the data with lower error and high accuracy rate. When extracting and plotting the importance of random forest in the code we can see the following:

	IncNodePurity
Blast.Furnace.Slag	11066.801
Water	24434.640
Coarse.Aggregate	10528.946
Age	50511.355
Cement	34792.586
Fly.Ash	9079.506
Superplasticizer	17187.622
Fine.Aggregate	12678.759

Based on the test results, the Random Forest model seems to be the most accurate predictor. A more thorough understanding of the variables affecting concrete strength may result from this knowledge, which might enhance decision-making and optimise the design and formulation of concrete mixes.

Conclusion

In conclusion, Among the five different models that was used in the regression problems, the most accurate one is the Random Forest as the model is consistent. In case of the regression model, MAE, MSE, EMSE, R2 and adjusted R2 we, R2 gave more variance between predictor variable and the target variable. The variance measured by R2 is the ratio of two different variables one being dependent variable which is predictable from the independent variable. Evaluation metrics help find or explain the variability in the data.

Bibliography

- Ibm.com. ² (2024). *What Is Linear Regression? | IBM*. [online] Available at: <https://www.ibm.com/topics/linear-regression#:~:text=IBM-,What%20is%20linear%20regression%3F,is%20called%20the%20independent%20variable.> [Accessed 10 May 2024].
- Ibm.com. (2024). *What Is Random Forest? | IBM*. [online] Available at: <https://www.ibm.com/topics/random-forest> [Accessed 10 May 2024].
- ⁸ GeeksforGeeks (2017). *Decision Tree*. [online] GeeksforGeeks. Available at: <https://www.geeksforgeeks.org/decision-tree/> [Accessed 10 May 2024].

AML Report part-2

ORIGINALITY REPORT

20%

SIMILARITY INDEX

16%

INTERNET SOURCES

7%

PUBLICATIONS

16%

STUDENT PAPERS

PRIMARY SOURCES

1	dergipark.org.tr Internet Source	4%
2	Submitted to University of Bradford Student Paper	3%
3	Submitted to University of Surrey Student Paper	2%
4	user-jbdbkvi.cld.bz Internet Source	2%
5	Submitted to Liverpool John Moores University Student Paper	2%
6	Submitted to University of North Texas Student Paper	2%
7	escholarship.org Internet Source	1%
8	Submitted to Leeds Beckett University Student Paper	1%
9	Submitted to University of Wales Institute, Cardiff	1%

10	repository.effatuniversity.edu.sa Internet Source	1 %
11	Submitted to University of Bolton Student Paper	1 %
12	cs229.stanford.edu Internet Source	1 %
13	Submitted to University of the Sunshine Coast Student Paper	<1 %
14	fastercapital.com Internet Source	<1 %
15	research.tue.nl Internet Source	<1 %
16	Yuhui Lyu, Madura Pathirage, Elham Ramyar, Wing Kam Liu, Gianluca Cusatis. "Machine learning meta-models for fast parameter identification of the lattice discrete particle model", Computational Mechanics, 2023 Publication	<1 %

Exclude quotes Off

Exclude matches Off

Exclude bibliography Off